

Assignment no :3

Name : CHO VATIYA JAY MAHESHBHAI

Student ID : 1161980

Course: COMP5421 Deep Learning

Topic: Datasets: CIFAR100

Requirements :

Last digit of ID: 0

Data set used: CIFAR100

Second last digit of ID: 8

DCNN used: Dense-Net

Notes : In this requirement 2 both convolutional neural network architectures such as Densnet and Alexnet i tried to implement in order to check the accuracy of the model which you can see below in the given output.

Explanation:

Q-Compare the performance gap between the pretrained DCNN model in the condition of transfer learning, and your customized model in assignment 2 in the condition of training from scratch. Explain why you obtained such results, and give a brief discussion about it.

Ans: Obviously, **the performance is better in transfer learning.** To elaborate it, In transfer learning, we have used the pretrained dense net model on the ImageNet which has 16 million data and directly use it to the cifar100.

So, the **testing time(93s/epoch)** is quite **less** compared to **training from scratch(137s/epoch)** and **accuracy(63.54)** is quite **higher** than **training from scratch(52.03)**. Hence, we can get good results from pretrained models rather than scratch we had done in the assignment 2 by using the customized model.

IMPLEMENTATION :

```
# import all the libraries
import numpy as np
import sklearn.metrics as metrics
from keras.applications import densenet
from keras.datasets import cifar100
from keras.utils import np_utils
from keras.optimizers import Adam
from keras.models import Sequential
from keras.layers import Dense, Flatten, Dropout
from keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt

# for removing warnings
import warnings
warnings.filterwarnings('ignore')

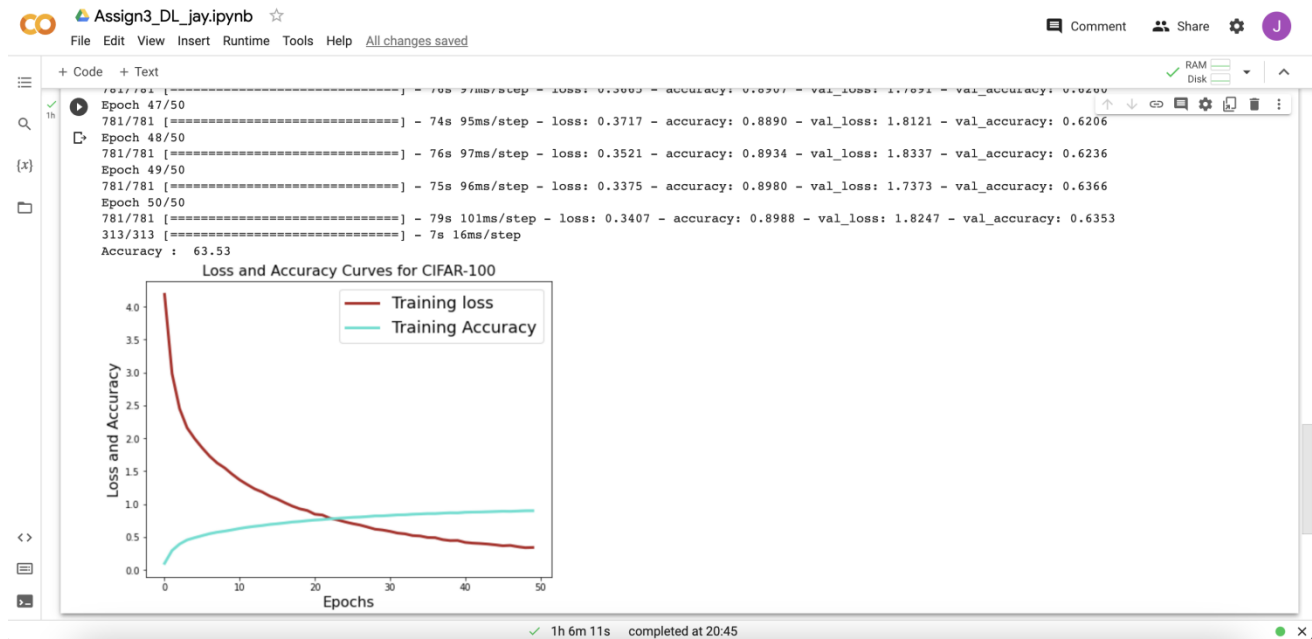
# define imagenet-pretrained model for densenet(cifar100)
model = Sequential()
model.add(densenet.DenseNet121(weights='imagenet', include_top=False,
input_shape=(32,32,3), pooling='max'))
# adding dense layer for flatten
model.add(Dense(256, activation='relu'))
# deactivating 50% nodes
model.add(Dropout(0.5))
model.add(Dense(100, activation='softmax'))
# find summary
model.summary()
# Splitting training and testing set
(cifarx_train, cifary_train), (cifarx_test, cifary_test) = cifar100.load_data()
# Converting to float
cifarx_train = cifarx_train.astype('float32')
cifarx_test = cifarx_test.astype('float32')
# converting data into normalize form
cifarx_train = densenet.preprocess_input(cifarx_train)
cifarx_test = densenet.preprocess_input(cifarx_test)
# data augmentation
datagen = ImageDataGenerator(rotation_range=15, width_shift_range=0.1, height_shift_range=0.1,
horizontal_flip=True)
datagen.fit(cifarx_train)
# one-hot encoding
Y_train = np_utils.to_categorical(cifary_train, 100)
Y_test = np_utils.to_categorical(cifary_test, 100)
# Using Adam optimizer to speed up training and set learning rate 0.001
optimizer = Adam(lr=1e-4)
# compile the model
model.compile(loss='categorical_crossentropy', optimizer=optimizer, metrics=["accuracy"])
# train the model
history = model.fit(datagen.flow(cifarx_train, Y_train, batch_size=64),
steps_per_epoch=len(cifarx_train) / 64, epochs=50,
```

```

validation_data=(cifarx_test, Y_test))
Preds = model.predict(cifarx_test)
y_Pred = np.argmax(Preds, axis=1)
y_true = cifary_test.flatten()
# finding accuracy and loss
accuracy = metrics.accuracy_score(y_true, y_Pred) * 100
# print testing accuracy
print("Accuracy : ", accuracy)
# Define plotchart function
def plotchart(history, value):
    plt.figure(figsize=[8,6])
    plt.plot(history.history['loss'], 'firebrick', linewidth=3.0)
    plt.plot(history.history['accuracy'], 'turquoise', linewidth=3.0)
    plt.legend(['Training loss', 'Training Accuracy'], fontsize=18)
    plt.xlabel('Epochs', fontsize=16)
    plt.ylabel('Loss and Accuracy', fontsize=16)
    plt.title('Loss and Accuracy Curves for {}'.format(value), fontsize=16)
    plt.show()
# Plot the training history
plotchart(history, 'CIFAR-100')

```

Outputs:1 Here i have attached the screenshot of CIFAR100 datasets output which you can see below.



Assign3_DL_jay.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Code + Text

```

Epoch 36/50
781/781 [=====] - 76s 97ms/step - loss: 0.4922 - accuracy: 0.8550 - val_loss: 1.6227 - val_accuracy: 0.6402
Epoch 37/50
781/781 [=====] - 75s 96ms/step - loss: 0.4892 - accuracy: 0.8553 - val_loss: 1.7180 - val_accuracy: 0.6240
Epoch 38/50
781/781 [=====] - 75s 96ms/step - loss: 0.4597 - accuracy: 0.8621 - val_loss: 1.6067 - val_accuracy: 0.6400
Epoch 39/50
781/781 [=====] - 82s 105ms/step - loss: 0.4452 - accuracy: 0.8671 - val_loss: 1.7158 - val_accuracy: 0.6287
Epoch 40/50
781/781 [=====] - 76s 97ms/step - loss: 0.4469 - accuracy: 0.8661 - val_loss: 1.6664 - val_accuracy: 0.6353
Epoch 41/50
781/781 [=====] - 75s 97ms/step - loss: 0.4176 - accuracy: 0.8751 - val_loss: 1.7079 - val_accuracy: 0.6400
Epoch 42/50
781/781 [=====] - 74s 95ms/step - loss: 0.4068 - accuracy: 0.8778 - val_loss: 1.7185 - val_accuracy: 0.6380
Epoch 43/50
781/781 [=====] - 74s 95ms/step - loss: 0.4010 - accuracy: 0.8792 - val_loss: 1.6611 - val_accuracy: 0.6429
Epoch 44/50
781/781 [=====] - 74s 95ms/step - loss: 0.3923 - accuracy: 0.8828 - val_loss: 1.7785 - val_accuracy: 0.6253
Epoch 45/50
781/781 [=====] - 76s 97ms/step - loss: 0.3805 - accuracy: 0.8865 - val_loss: 1.7231 - val_accuracy: 0.6385
Epoch 46/50
781/781 [=====] - 76s 97ms/step - loss: 0.3665 - accuracy: 0.8907 - val_loss: 1.7891 - val_accuracy: 0.6260
Epoch 47/50
781/781 [=====] - 74s 95ms/step - loss: 0.3717 - accuracy: 0.8890 - val_loss: 1.8121 - val_accuracy: 0.6206
Epoch 48/50
781/781 [=====] - 76s 97ms/step - loss: 0.3521 - accuracy: 0.8934 - val_loss: 1.8337 - val_accuracy: 0.6236
Epoch 49/50
781/781 [=====] - 75s 96ms/step - loss: 0.3375 - accuracy: 0.8980 - val_loss: 1.7373 - val_accuracy: 0.6366
Epoch 50/50
781/781 [=====] - 79s 101ms/step - loss: 0.3407 - accuracy: 0.8988 - val_loss: 1.8247 - val_accuracy: 0.6353
313/313 [=====] - 7s 16ms/step
Accuracy : 63.53

```

Outputs:1.1

Assign3_DL_JAY.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Code + Text

```

Epoch 6/20
390/390 [=====] - 49s 124ms/step - loss: 1.7767 - accuracy: 0.5381 - val_loss: 1.6854 - val_accuracy: 0.5560
Epoch 7/20
390/390 [=====] - 49s 125ms/step - loss: 1.6344 - accuracy: 0.5723 - val_loss: 1.6727 - val_accuracy: 0.5530
Epoch 8/20
390/390 [=====] - 49s 125ms/step - loss: 1.5150 - accuracy: 0.5977 - val_loss: 1.5496 - val_accuracy: 0.5869
Epoch 9/20
390/390 [=====] - 49s 124ms/step - loss: 1.4187 - accuracy: 0.6212 - val_loss: 1.5379 - val_accuracy: 0.5904
Epoch 10/20
390/390 [=====] - 49s 124ms/step - loss: 1.3336 - accuracy: 0.6407 - val_loss: 1.5492 - val_accuracy: 0.5864
Epoch 11/20
390/390 [=====] - 54s 138ms/step - loss: 1.2618 - accuracy: 0.6603 - val_loss: 1.5803 - val_accuracy: 0.5880
Epoch 12/20
390/390 [=====] - 49s 126ms/step - loss: 1.1871 - accuracy: 0.6780 - val_loss: 1.5185 - val_accuracy: 0.6007
Epoch 13/20
390/390 [=====] - 49s 126ms/step - loss: 1.1214 - accuracy: 0.6933 - val_loss: 1.4781 - val_accuracy: 0.6140
Epoch 14/20
390/390 [=====] - 49s 126ms/step - loss: 1.0603 - accuracy: 0.7085 - val_loss: 1.4723 - val_accuracy: 0.6122
Epoch 15/20
390/390 [=====] - 48s 123ms/step - loss: 1.0098 - accuracy: 0.7227 - val_loss: 1.4822 - val_accuracy: 0.6167
Epoch 16/20
390/390 [=====] - 49s 124ms/step - loss: 0.9569 - accuracy: 0.7350 - val_loss: 1.5061 - val_accuracy: 0.6163
Epoch 17/20
390/390 [=====] - 48s 123ms/step - loss: 0.9006 - accuracy: 0.7484 - val_loss: 1.4691 - val_accuracy: 0.6243
Epoch 18/20
390/390 [=====] - 48s 122ms/step - loss: 0.8501 - accuracy: 0.7603 - val_loss: 1.5381 - val_accuracy: 0.6195
Epoch 19/20
390/390 [=====] - 48s 123ms/step - loss: 0.8103 - accuracy: 0.7694 - val_loss: 1.4607 - val_accuracy: 0.6355
Epoch 20/20
390/390 [=====] - 47s 120ms/step - loss: 0.7732 - accuracy: 0.7813 - val_loss: 1.4869 - val_accuracy: 0.6293
313/313 [=====] - 7s 16ms/step
Accuracy : 62.93

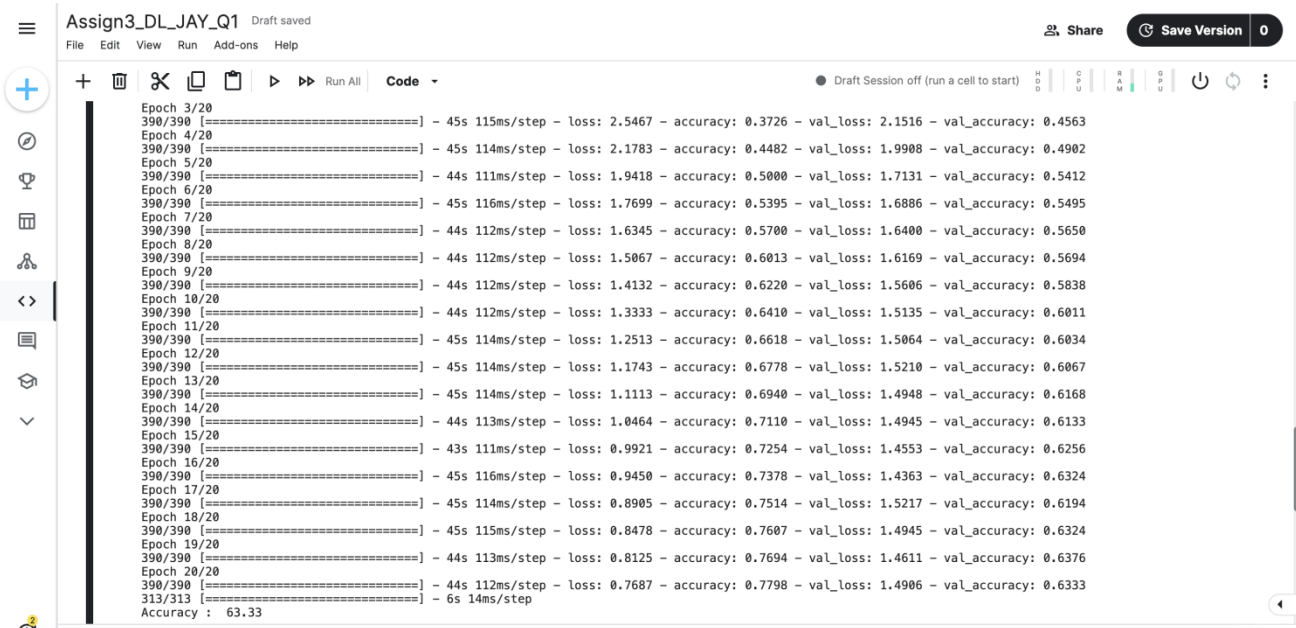
```

16m 52s completed at 00:11

Outputs:1.2



Outputs:1.3



Notes : In this question both convolutional neural network architectures such as Densnet and Alexnet i tried to implement in order to check the accuracy of the model which you can see below in the given output.

Requirement 2 :Use Freezing layers pretrained DCNN.

Q-Obtain the transfer learning features/deep features from a raw dataset.

```
#importing library
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms

# Declaring hyperparameters
epoch_no = 20 ; no_classes = 100
batch_size = 128 ; learning_rate = 0.1

# Define transformations for training and testing datasets
transform_train = transforms.Compose([transforms.RandomCrop(32,
padding=4),transforms.RandomHorizontalFlip(),transforms.ToTensor(),transforms.Nor
malize((0.5071, 0.4865, 0.4409), (0.2673, 0.2564, 0.2762))])
transform_test =
transforms.Compose([transforms.ToTensor(),transforms.Normalize((0.5071, 0.4865,
0.4409), (0.2673, 0.2564, 0.2762))])

# Loading CIFAR-100 dataset
cifer_train = torchvision.datasets.CIFAR100(root='./data', train=True, download=True,
transform=transform_train)
cifer_test = torchvision.datasets.CIFAR100(root='./data', train=False, download=True,
transform=transform_test)

# Defining data loaders for training and testing datasets
cifer_loader_train = torch.utils.data.DataLoader(dataset=cifer_train,
batch_size=batch_size, shuffle=True)
cifer_loader_test = torch.utils.data.DataLoader(dataset=cifer_test,
batch_size=batch_size, shuffle=False)

# Defining DenseNet-121 model and loss function and optimizer
model = torchvision.models.densenet121(pretrained=False, no_classes=no_classes)
```

```
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate, momentum=0.9,
weight_decay=1e-4)

# Training the model
total_step = len(cifer_loader_train)
for epoch in range(epoch_no):
    for i, (images, labels) in enumerate(cifer_loader_train):
        outputs = model(images)
        loss = criterion(outputs, labels)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        if (i+1) % 100 == 0:
            print('Epoch [{}/{}], Step [{}/{}], Loss: {:.4f}'
                .format(epoch+1, epoch_no, i+1, total_step, loss.item()))
# Testing the model
model.eval()
with torch.no_grad():
    correct = 0
    total = 0
    for images, labels in cifer_loader_test:
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
    accuracy = 100 * correct / total
    print('Test Accuracy of the model on the {} test images: {:.2f}%'.format(total, accuracy))
```

Outputs:2.1 Here i have attached the screenshot of obtain the transfer learning features from a raw dataset output which you can see below.

```

Users > jaychovatiya > Desktop > Assign3_DL_JAY > Assign3_DL_Q2.ipynb > import torch
+ Code + Markdown | Outline ...
Epoch [1/20], Step [100/391], Loss: 3.9687
Epoch [1/20], Step [300/391], Loss: 3.8321
Test Accuracy of the model on the 10000 test images: 14.52%
Epoch [2/20], Step [100/391], Loss: 3.6442
Epoch [2/20], Step [200/391], Loss: 3.5376
Epoch [2/20], Step [300/391], Loss: 3.3820
Test Accuracy of the model on the 10000 test images: 18.20%
Epoch [3/20], Step [100/391], Loss: 3.3561
Epoch [3/20], Step [200/391], Loss: 3.4391
Epoch [3/20], Step [300/391], Loss: 3.1639
Test Accuracy of the model on the 10000 test images: 22.31%
Epoch [4/20], Step [100/391], Loss: 3.3098
Epoch [4/20], Step [200/391], Loss: 3.0604
Epoch [4/20], Step [300/391], Loss: 2.8664
Test Accuracy of the model on the 10000 test images: 24.24%
Epoch [5/20], Step [100/391], Loss: 2.7170
Epoch [5/20], Step [200/391], Loss: 2.8574
Epoch [5/20], Step [300/391], Loss: 2.8250
Test Accuracy of the model on the 10000 test images: 28.37%
Epoch [6/20], Step [100/391], Loss: 3.0163
Epoch [6/20], Step [200/391], Loss: 2.6568
Epoch [6/20], Step [300/391], Loss: 2.5661
...
Epoch [20/20], Step [100/391], Loss: 1.9101
Epoch [20/20], Step [200/391], Loss: 1.9694
Epoch [20/20], Step [300/391], Loss: 2.0909
Test Accuracy of the model on the 10000 test images: 42.58%

```

Outputs:2.2



Assign3_DL_Q2 Draft saved

File Edit View Run Add-ons Help

Share Save Version 2

Run All Code

Draft Session off (run a cell to start)

```

Test Accuracy of the model on the 10000 test images: 48.64%
Epoch [42/50], Step [100/391], Loss: 1.3616
Epoch [42/50], Step [200/391], Loss: 1.5264
Epoch [42/50], Step [300/391], Loss: 1.1891
Test Accuracy of the model on the 10000 test images: 48.11%
Epoch [43/50], Step [100/391], Loss: 0.9936
Epoch [43/50], Step [200/391], Loss: 1.2273
Epoch [43/50], Step [300/391], Loss: 1.3180
Test Accuracy of the model on the 10000 test images: 47.18%
Epoch [44/50], Step [100/391], Loss: 1.0486
Epoch [44/50], Step [200/391], Loss: 1.1764
Epoch [44/50], Step [300/391], Loss: 1.1359
Test Accuracy of the model on the 10000 test images: 48.54%
Epoch [45/50], Step [100/391], Loss: 1.1811
Epoch [45/50], Step [200/391], Loss: 1.2310
Epoch [45/50], Step [300/391], Loss: 1.2390
Test Accuracy of the model on the 10000 test images: 47.25%
Epoch [46/50], Step [100/391], Loss: 1.2571
Epoch [46/50], Step [200/391], Loss: 1.2220
Epoch [46/50], Step [300/391], Loss: 1.3411
Test Accuracy of the model on the 10000 test images: 48.68%
Epoch [47/50], Step [100/391], Loss: 1.2593
Epoch [47/50], Step [200/391], Loss: 1.4388
Epoch [47/50], Step [300/391], Loss: 1.2230
Test Accuracy of the model on the 10000 test images: 48.49%
Epoch [48/50], Step [100/391], Loss: 1.2733
Epoch [48/50], Step [200/391], Loss: 1.2387
Epoch [48/50], Step [300/391], Loss: 1.0115
Test Accuracy of the model on the 10000 test images: 49.68%
Epoch [49/50], Step [100/391], Loss: 1.0856
Epoch [49/50], Step [200/391], Loss: 0.9243
Epoch [49/50], Step [300/391], Loss: 1.2783
Test Accuracy of the model on the 10000 test images: 48.58%
Epoch [50/50], Step [100/391], Loss: 1.0176
Epoch [50/50], Step [200/391], Loss: 0.7986
Epoch [50/50], Step [300/391], Loss: 1.0023
Test Accuracy of the model on the 10000 test images: 47.52%

```

Outputs:2.3

Users > jaychovatiya > Desktop > Assign3_DL_JAY > Alexnet.ipynb > import torch


Search Markdown Outline

Select Kernel

```

Epoch [1/15], Step [200/391], Loss: 4.6043
Epoch [1/15], Step [300/391], Loss: 4.5378
Test Accuracy of the model on the 10000 test images: 3.83%
Epoch [2/15], Step [100/391], Loss: 4.1526
Epoch [2/15], Step [200/391], Loss: 4.0647
Epoch [2/15], Step [300/391], Loss: 3.8441
Test Accuracy of the model on the 10000 test images: 12.93%
Epoch [3/15], Step [100/391], Loss: 3.5918
Epoch [3/15], Step [200/391], Loss: 3.3421
Epoch [3/15], Step [300/391], Loss: 3.6319
Test Accuracy of the model on the 10000 test images: 22.27%
Epoch [4/15], Step [100/391], Loss: 3.2504
Epoch [4/15], Step [200/391], Loss: 2.9298
Epoch [4/15], Step [300/391], Loss: 3.1729
Test Accuracy of the model on the 10000 test images: 27.51%
Epoch [5/15], Step [100/391], Loss: 2.7019
Epoch [5/15], Step [200/391], Loss: 2.8515
Epoch [5/15], Step [300/391], Loss: 2.5224
Test Accuracy of the model on the 10000 test images: 33.45%
Epoch [6/15], Step [100/391], Loss: 2.5199
Epoch [6/15], Step [200/391], Loss: 2.3374
Epoch [6/15], Step [300/391], Loss: 2.3628
...
Epoch [15/15], Step [100/391], Loss: 0.7801
Epoch [15/15], Step [200/391], Loss: 0.6985
Epoch [15/15], Step [300/391], Loss: 0.4889
Test Accuracy of the model on the 10000 test images: 48.92%

```



I have tried to implement using the keras as well in which i received the output like this.

ASSIGNMENT3_DL_JAY(2) Draft saved

File Edit View Run Add-ons Help

Download data from https://storage.googleapis.com/tensorflow/keras-applications/densenet/densenet121_weights_tf_dim_ordering_tf_kernels_notop.h5
29084464/29084464 [=====] - 0s 0us/step
Model: "sequential"

Layer (type)	Output Shape	Param #
densenet121 (Functional)	(None, 1024)	7037504
dense (Dense)	(None, 256)	262400
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 100)	25700

Total params: 7,325,604
Trainable params: 616,228
Non-trainable params: 6,709,376

Download data from <https://www.cs.toronto.edu/~kriz/cifar-100-python.tar.gz>
169001437/169001437 [=====] - 2s 0us/step
1563/1563 [=====] - 36s 17ms/step
313/313 [=====] - 5s 16ms/step

Epoch: 1 Accuracy: 2.19
Epoch: 2 Accuracy: 2.19
Epoch: 3 Accuracy: 2.19
Epoch: 4 Accuracy: 2.19
Epoch: 5 Accuracy: 2.19
Epoch: 6 Accuracy: 2.19
Epoch: 7 Accuracy: 2.19
Epoch: 8 Accuracy: 2.19
Epoch: 9 Accuracy: 2.19
Epoch: 10 Accuracy: 2.19
Accuracy : 2.19
Shape of deep features (train): (50000, 1024)
Shape of deep features (test): (10000, 1024)